EN3160 Assignment 3: Neural Networks

Name: Balasooriya B.R.B.D.
Index number: 230070T
GitHub repository: https://github.com/banula33/A03

```
import os
os.environ["TORCH_HOME"] = "./data/torch"
import random
import time
import json
from pathlib import Path
import numpy as np
import matplotlib.pyplot as plt
import torch
from torch import nn
from torch.utils.data import DataLoader, TensorDataset
from torchvision import datasets, transforms, models
from IPython.display import display, Markdown

SEED = 42
def set_seed():
    random.seed(SEED)
    np.random.seed(SEED)
    torch.manual_seed(SEED)
    if torch.cuda.is_available():
        torch.cuda.manual_seed_all(SEED)
    torch.backends.cudnn.deterministic = True
    torch.backends.cudnn.benchmark = False

set_seed()
torch.set_num_threads(4)
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
Path("results").mkdir(exist_ok=True)
print("Device:", device)
print("PyTorch:", torch.__version__)
print("Seed:", SEED)

def make_loader(dataset, batch_size, shuffle=False):
    generator = torch.Generator().manual_seed(SEED)
    return DataLoader(dataset, batch_size=batch_size, shuffle=shuffle,
                      num_workers=0, generator=generator)

def save_figure(name):
    plt.tight_layout()
    plt.savefig("results/" + name + ".png", dpi=150, bbox_inches="tight")
    plt.show()

def plot_history(history, title, name, evaluation="Test"):
    epochs = range(1, len(history) + 1)
    fig, axes = plt.subplots(1, 2, figsize=(10, 3.5))
    for ax, metric in zip(axes, ["loss", "acc"]):
        ax.plot(epochs, [h["train_" + metric] for h in history], "o-", label="Train")
        ax.plot(epochs, [h["test_" + metric] for h in history], "o-", label=evaluation)
        ax.set(xlabel="Epoch", ylabel="Loss" if metric == "loss" else "Accuracy (%)")
        ax.grid(alpha=0.3)
        ax.legend()
    fig.suptitle(title)
    save_figure(name)

def print_history(history):
    print("Epoch  Train loss  Eval loss  Train acc  Eval acc")
    for i, h in enumerate(history, 1):
        print(f"{i:5d}  {h['train_loss']:10.4f}  {h['test_loss']:9.4f}"
              f"  {h['train_acc']:9.2f}  {h['test_acc']:8.2f}")
```

```
Device: cpu
PyTorch: 2.14.1+cpu
Seed: 42
```

1. CIFAR-10 with manual backpropagation

The network is 3072 -> 100 sigmoid units -> 10 output logits. Softmax turns the logits into class probabilities. Cross-entropy is calculated with log-softmax for numerical stability. For a batch of N images, the output gradient is `(probabilities - one_hot_labels) / N`. The hidden-layer gradient includes the sigmoid derivative `h * (1 - h)`.

The batch size is 50, as in the supplied listing. The learning rate starts at 0.1 and is multiplied by 0.9 after each epoch. Xavier-style initialization replaces the very small initialization in the listing; the learning rate is also adjusted for mean cross-entropy. No optimizer or automatic differentiation is used to train this network.

```
cifar_train = datasets.CIFAR10("data", train=True, download=True)
cifar_test = datasets.CIFAR10("data", train=False, download=True)

def cifar_tensors(dataset):
    x = torch.tensor(dataset.data).permute(0, 3, 1, 2).float() / 255
    x = ((x - 0.5) / 0.5).reshape(-1, 3072)
    return TensorDataset(x, torch.tensor(dataset.targets))

cifar_train_data = cifar_tensors(cifar_train)
cifar_test_data = cifar_tensors(cifar_test)
cifar_train_loader = make_loader(cifar_train_data, 50, True)
cifar_train_eval = make_loader(cifar_train_data, 256)
cifar_test_loader = make_loader(cifar_test_data, 256)
print("CIFAR-10 train/test:", len(cifar_train_data), len(cifar_test_data))

set_seed()
W1 = torch.randn(3072, 100, device=device) * np.sqrt(2 / (3072 + 100))
b1 = torch.zeros(100, device=device)
W2 = torch.randn(100, 10, device=device) * np.sqrt(2 / (100 + 10))
b2 = torch.zeros(10, device=device)

def dense_forward(x):
    hidden = torch.sigmoid(x @ W1 + b1)
    return hidden, hidden @ W2 + b2

@torch.no_grad()
def evaluate_dense(loader):
    loss_sum, correct, total = 0.0, 0, 0
    for x, y in loader:
        x, y = x.to(device), y.to(device)
        _, logits = dense_forward(x)
        loss_sum += nn.functional.cross_entropy(logits, y, reduction="sum").item()
        correct += (logits.argmax(1) == y).sum().item()
        total += len(y)
    return loss_sum / total, 100 * correct / total
```

```
CIFAR-10 train/test: 50000 10000
```

```
dense_history = []
learning_rate = 0.1
start = time.perf_counter()
for epoch in range(10):
    for x, y in cifar_train_loader:
        x, y = x.to(device), y.to(device)
        hidden, logits = dense_forward(x)
        log_prob = torch.log_softmax(logits, dim=1)
        loss = -log_prob[torch.arange(len(y), device=device), y].mean()
        targets = nn.functional.one_hot(y, 10).float()

        dz2 = (log_prob.exp() - targets) / len(y)
        dW2 = hidden.T @ dz2
        db2 = dz2.sum(0)
        dz1 = (dz2 @ W2.T) * hidden * (1 - hidden)
        dW1 = x.T @ dz1
        db1 = dz1.sum(0)

        W1 -= learning_rate * dW1
        b1 -= learning_rate * db1
        W2 -= learning_rate * dW2
        b2 -= learning_rate * db2
```

```
    train_loss, train_acc = evaluate_dense(cifar_train_eval)
    test_loss, test_acc = evaluate_dense(cifar_test_loader)
    dense_history.append(dict(train_loss=train_loss, train_acc=train_acc,
                              test_loss=test_loss, test_acc=test_acc))
    print(f"Epoch {epoch + 1}: train {train_acc:.2f}%, test {test_acc:.2f}%", flush=True)
    learning_rate *= 0.9
dense_seconds = time.perf_counter() - start
print_history(dense_history)
plot_history(dense_history, "CIFAR-10: manual dense network", "cifar_curves")
```

```
Epoch 1: train 41.67%, test 41.50%

Epoch 2: train 44.43%, test 42.59%

Epoch 3: train 46.45%, test 44.02%

Epoch 4: train 47.79%, test 45.23%

Epoch 5: train 49.63%, test 46.03%

Epoch 6: train 50.66%, test 46.56%

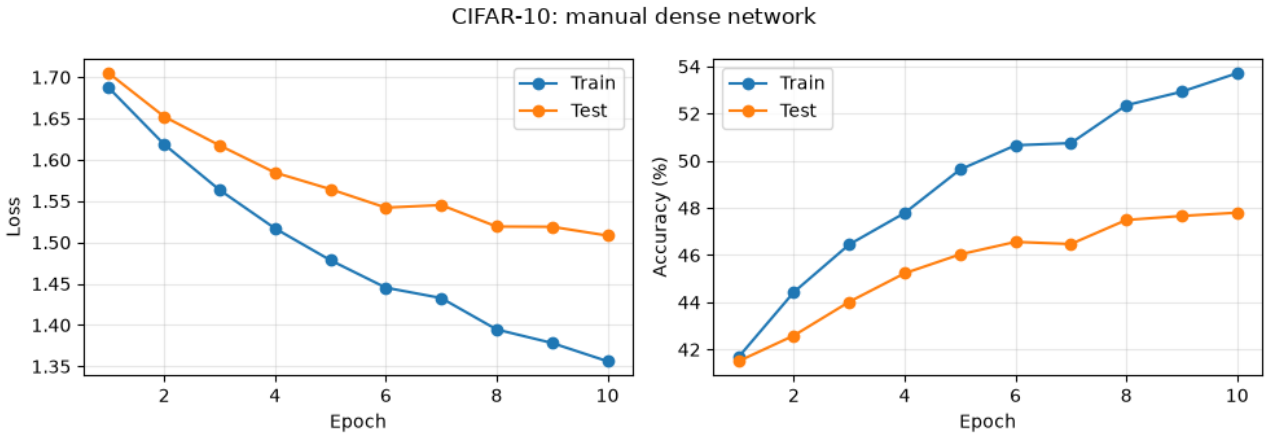
Epoch 7: train 50.75%, test 46.47%

Epoch 8: train 52.35%, test 47.49%

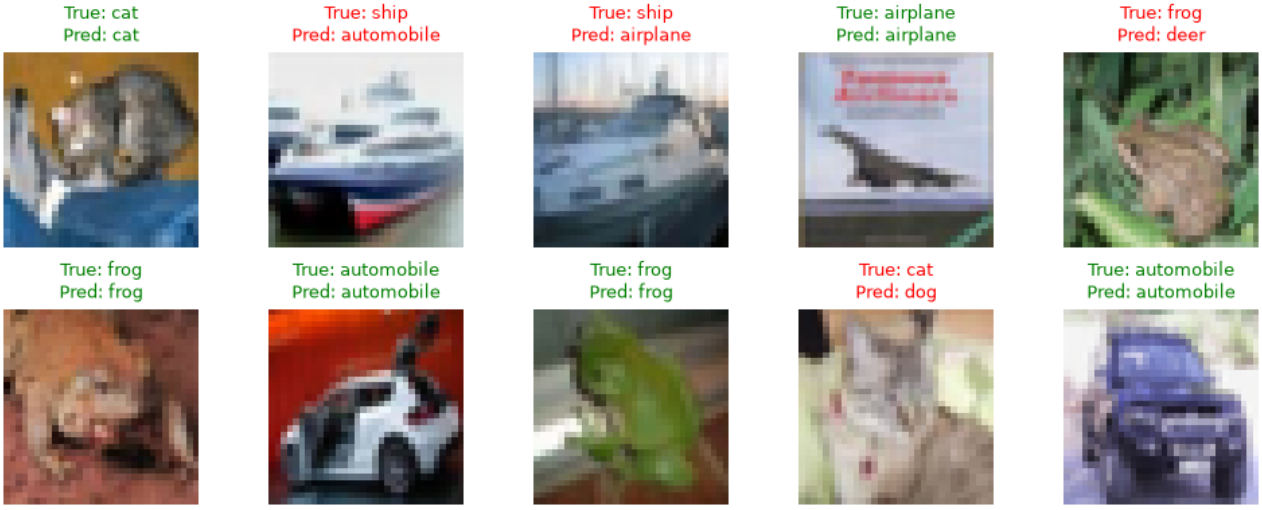
Epoch 9: train 52.93%, test 47.66%

Epoch 10: train 53.71%, test 47.80%
```

Epoch	Train loss	Eval loss	Train acc	Eval acc
1	1.6871	1.7048	41.67	41.50
2	1.6186	1.6522	44.43	42.59
3	1.5635	1.6174	46.45	44.02
4	1.5169	1.5845	47.79	45.23
5	1.4784	1.5644	49.63	46.03
6	1.4452	1.5421	50.66	46.56
7	1.4326	1.5452	50.75	46.47
8	1.3946	1.5192	52.35	47.49
9	1.3782	1.5189	52.93	47.66
10	1.3561	1.5083	53.71	47.80




```
fig, axes = plt.subplots(2, 5, figsize=(10, 4))
with torch.no_grad():
    _, logits = dense_forward(cifar_test_data.tensors[0][:10].to(device))
predictions = logits.argmax(1).cpu().tolist()
for i, ax in enumerate(axes.flat):
    true_label = cifar_test.targets[i]
    ax.imshow(cifar_test.data[i])
    ax.set_title(f"True: {cifar_test.classes[true_label]}\nPred:
{cifar_test.classes[predictions[i]]}",
                 fontsize=9, color="green" if predictions[i] == true_label else "red")
    ax.axis("off")
save_figure("cifar_predictions")
h = dense_history[-1]
```



Discussion

After 10 epochs, training accuracy is **53.71%** and test accuracy is **47.80%**. The gap is **5.91 percentage points**. Test loss changes from 1.7048 in epoch 1 to 1.5083 in epoch 10. Training accuracy indicates how well the model fits the training set; test accuracy measures performance on unseen images. The gap indicates some overfitting if training accuracy is higher. Flattening removes the explicit spatial organization of the image, and 100 sigmoid units have limited capacity for varied natural images. The example predictions show individual successes and mistakes; they are the first ten test images, not a selection of only correct predictions.

2. MNIST: LeNet-5 and a small modern CNN

Both models use the same 60,000 training images and 10,000 test images, normalization, batch size 128, Adam optimizer (learning rate 0.001), cross-entropy loss and 10 epochs. The seed and shuffled batch sequence are reset for each model. No data augmentation is used. The test set is monitored as requested, but is not used for tuning or checkpoint selection.

LeNet-5 uses 6 and 16 convolutional channels, tanh activations, average pooling and fully connected layers of 120, 84 and 10 units. Padding the first convolution allows 28x28 MNIST images. This is the usual fully connected PyTorch adaptation: the original LeNet-5's sparse connections and RBF output are replaced by dense connections and logits. The modern model uses three convolutions (16, 32 and 64 channels), batch normalization, ReLU activations, max pooling and dropout before the final classifier.

```
mnist_train = datasets.MNIST("data", train=True, download=True)
mnist_test = datasets.MNIST("data", train=False, download=True)
def mnist_tensors(dataset):
    x = dataset.data.unsqueeze(1).float() / 255
    return TensorDataset((x - 0.1307) / 0.3081, dataset.targets)
mnist_train_data = mnist_tensors(mnist_train)
mnist_test_data = mnist_tensors(mnist_test)
mnist_train_eval = make_loader(mnist_train_data, 256)
mnist_test_loader = make_loader(mnist_test_data, 256)
print("MNIST train/test:", len(mnist_train_data), len(mnist_test_data))

class LeNet5(nn.Module):
    def __init__(self):
        super().__init__()
        self.layers = nn.Sequential(
            nn.Conv2d(1, 6, 5, padding=2), nn.Tanh(), nn.AvgPool2d(2),
            nn.Conv2d(6, 16, 5), nn.Tanh(), nn.AvgPool2d(2),
            nn.Flatten(), nn.Linear(16 * 5 * 5, 120), nn.Tanh(),
            nn.Linear(120, 84), nn.Tanh(), nn.Linear(84, 10))
    def forward(self, x):
        return self.layers(x)

class SmallCNN(nn.Module):
    def __init__(self):
```

```
        super().__init__()
        self.layers = nn.Sequential(
            nn.Conv2d(1, 16, 3, padding=1), nn.BatchNorm2d(16),
            nn.ReLU(), nn.MaxPool2d(2),
            nn.Conv2d(16, 32, 3, padding=1), nn.BatchNorm2d(32),
            nn.ReLU(), nn.MaxPool2d(2),
            nn.Conv2d(32, 64, 3, padding=1), nn.BatchNorm2d(64),
            nn.ReLU(), nn.MaxPool2d(2),
            nn.Flatten(), nn.Dropout(0.5), nn.Linear(64 * 3 * 3, 10))
    def forward(self, x):
        return self.layers(x)

@torch.no_grad()
def evaluate(model, loader):
    model.eval()
    loss_sum, correct, total = 0.0, 0, 0
    labels, predictions = [], []
    for x, y in loader:
        x, y = x.to(device), y.to(device)
        logits = model(x)
        loss_sum += nn.functional.cross_entropy(logits, y, reduction="sum").item()
        pred = logits.argmax(1)
        correct += (pred == y).sum().item()
        total += len(y)
        labels.extend(y.cpu().tolist())
        predictions.extend(pred.cpu().tolist())
    return loss_sum / total, 100 * correct / total, np.array(labels), np.array(predictions)

def train_model(model, train_loader, train_eval, test_loader, optimizer,
                epochs=10, scheduler=None, frozen_backbone=False):
    history = []
    start = time.perf_counter()
    for epoch in range(epochs):
        model.train()
        if frozen_backbone:
            model.eval()  # Freeze batch-normalization statistics as well as weights.
            model.fc.train()
        loss_sum = 0.0
        for x, y in train_loader:
            x, y = x.to(device), y.to(device)
            optimizer.zero_grad()
            logits = model(x)
            loss = nn.functional.cross_entropy(logits, y)
            loss.backward()
            optimizer.step()
            loss_sum += loss.item() * len(y)
        train_loss, train_acc, _, _ = evaluate(model, train_eval)
        test_loss, test_acc, _, _ = evaluate(model, test_loader)
        history.append(dict(train_loss=train_loss, train_acc=train_acc,
                            test_loss=test_loss, test_acc=test_acc,
                            online_loss=loss_sum / len(train_loader.dataset)))
        print(f"Epoch {epoch + 1}: train {train_acc:.2f}%, eval {test_acc:.2f}%", flush=True)
        if scheduler is not None:
            scheduler.step()
    return history, time.perf_counter() - start

def confusion_plot(labels, predictions, names, title, filename):
    matrix = np.zeros((len(names), len(names)), dtype=int)
    np.add.at(matrix, (labels, predictions), 1)
    fig, ax = plt.subplots(figsize=(6, 5))
    image = ax.imshow(matrix, cmap="Blues")
    fig.colorbar(image, ax=ax)
    for i in range(len(names)):
        for j in range(len(names)):
            ax.text(j, i, str(matrix[i, j]), ha="center", va="center", fontsize=8,
                    color="white" if matrix[i, j] > matrix.max() / 2 else "black")
    ax.set(xticks=range(len(names)), yticks=range(len(names)),
           xticklabels=names, yticklabels=names, xlabel="Predicted class",
           ylabel="True class", title=title)
    save_figure(filename)
    return matrix
```

```
MNIST train/test: 60000 10000
```

```
mnist_results = {}
mnist_models = {}
for name, model_class in [("LeNet-5", LeNet5), ("Small CNN", SmallCNN)]:
    set_seed()
    model = model_class().to(device)
    train_loader = make_loader(mnist_train_data, 128, True)
    optimizer = torch.optim.Adam(model.parameters(), lr=0.001)
    print("\n" + name, flush=True)
    history, seconds = train_model(model, train_loader, mnist_train_eval,
                                   mnist_test_loader, optimizer)
    _, _, labels, predictions = evaluate(model, mnist_test_loader)
    mnist_results[name] = dict(history=history, seconds=seconds,
                               parameters=sum(p.numel() for p in model.parameters()))
    mnist_models[name] = model
    print_history(history)
    plot_history(history, name + ": MNIST", name.replace(" ", "_") + "_curves")
    matrix = confusion_plot(labels, predictions, list(range(10)),
                            name + ": test confusion matrix", name.replace(" ", "_") + "_confusion")
    mnist_results[name]["confusion"] = matrix.tolist()
```

```
LeNet-5

Epoch 1: train 96.72%, eval 96.76%

Epoch 2: train 98.11%, eval 97.90%

Epoch 3: train 98.62%, eval 98.28%

Epoch 4: train 98.73%, eval 98.37%

Epoch 5: train 99.29%, eval 98.75%

Epoch 6: train 99.18%, eval 98.77%

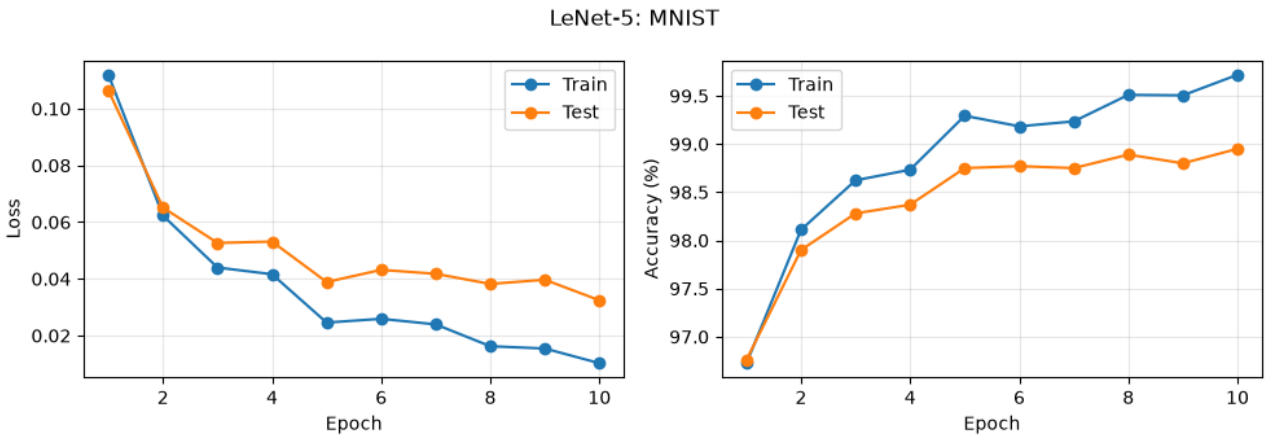
Epoch 7: train 99.23%, eval 98.75%

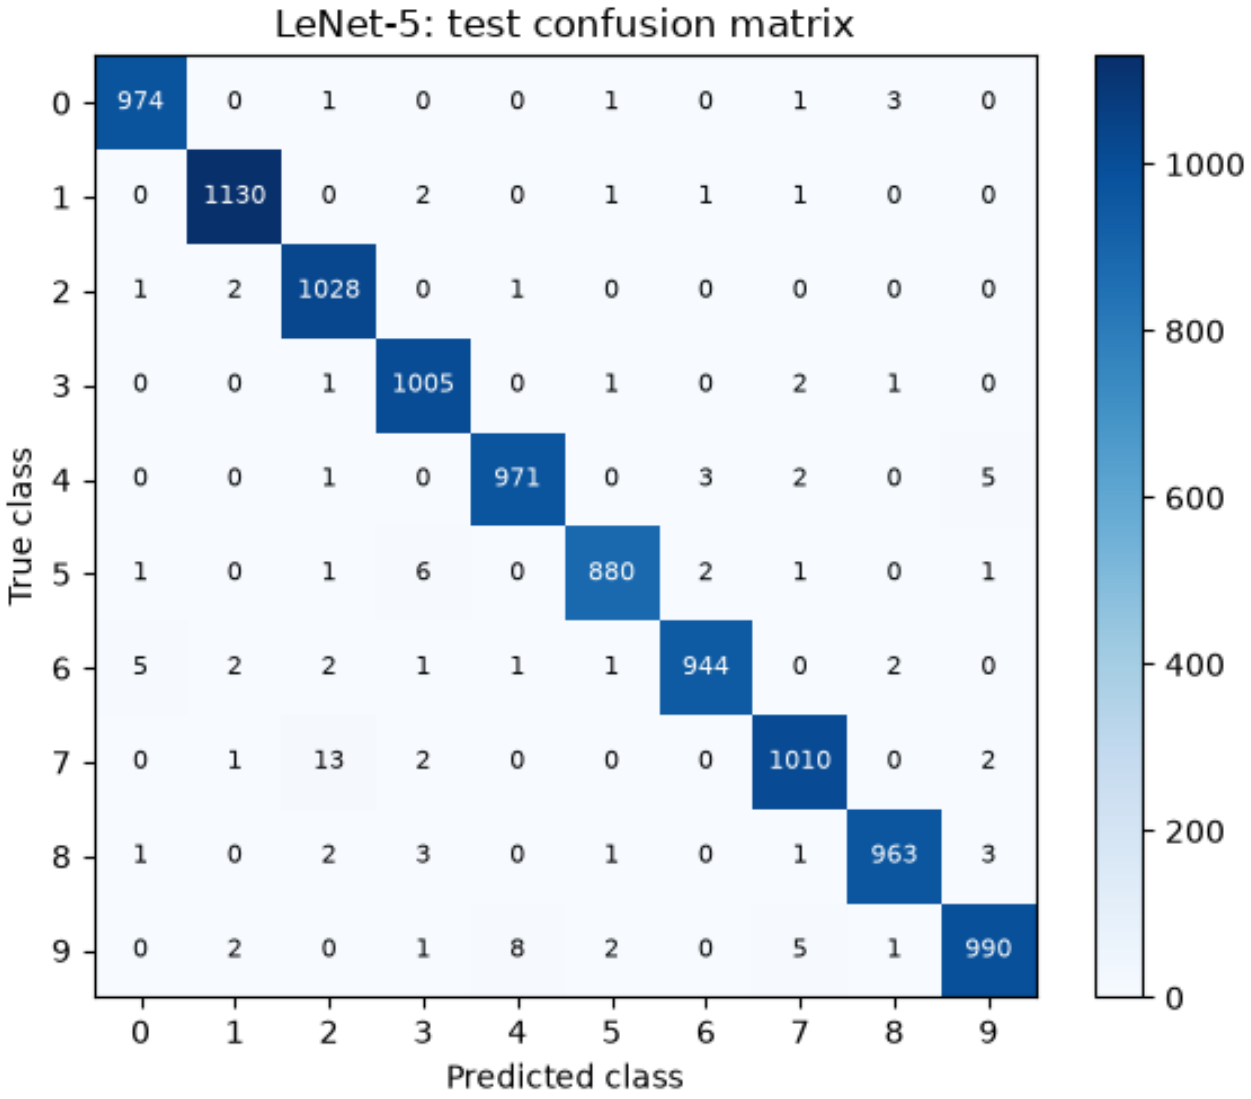
Epoch 8: train 99.51%, eval 98.89%

Epoch 9: train 99.50%, eval 98.80%

Epoch 10: train 99.72%, eval 98.95%
```

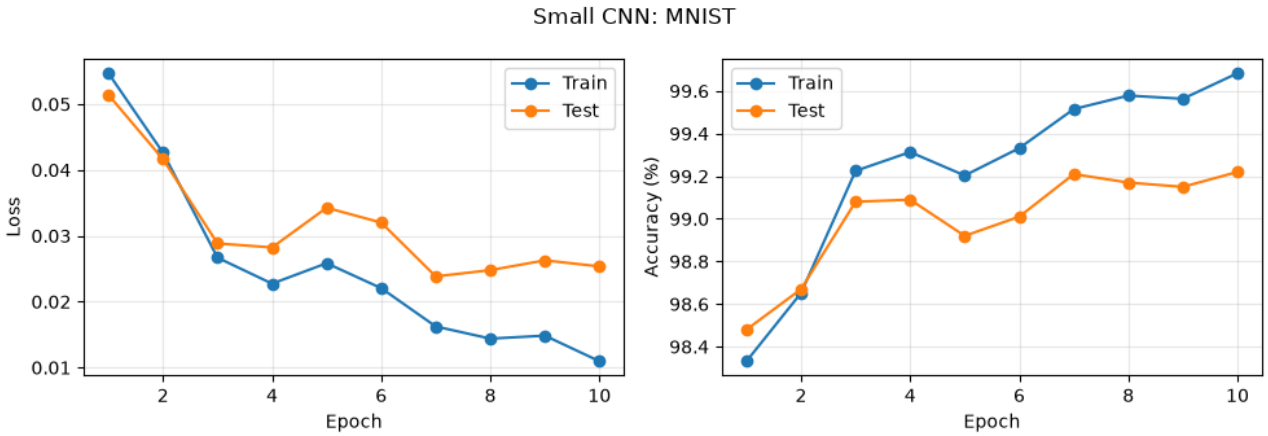
Epoch	Train loss	Eval loss	Train acc	Eval acc
1	0.1120	0.1065	96.72	96.76
2	0.0623	0.0651	98.11	97.90
3	0.0440	0.0527	98.62	98.28
4	0.0416	0.0531	98.73	98.37
5	0.0246	0.0389	99.29	98.75
6	0.0259	0.0432	99.18	98.77
7	0.0239	0.0418	99.23	98.75
8	0.0162	0.0382	99.51	98.89
9	0.0154	0.0397	99.50	98.80
10	0.0102	0.0324	99.72	98.95

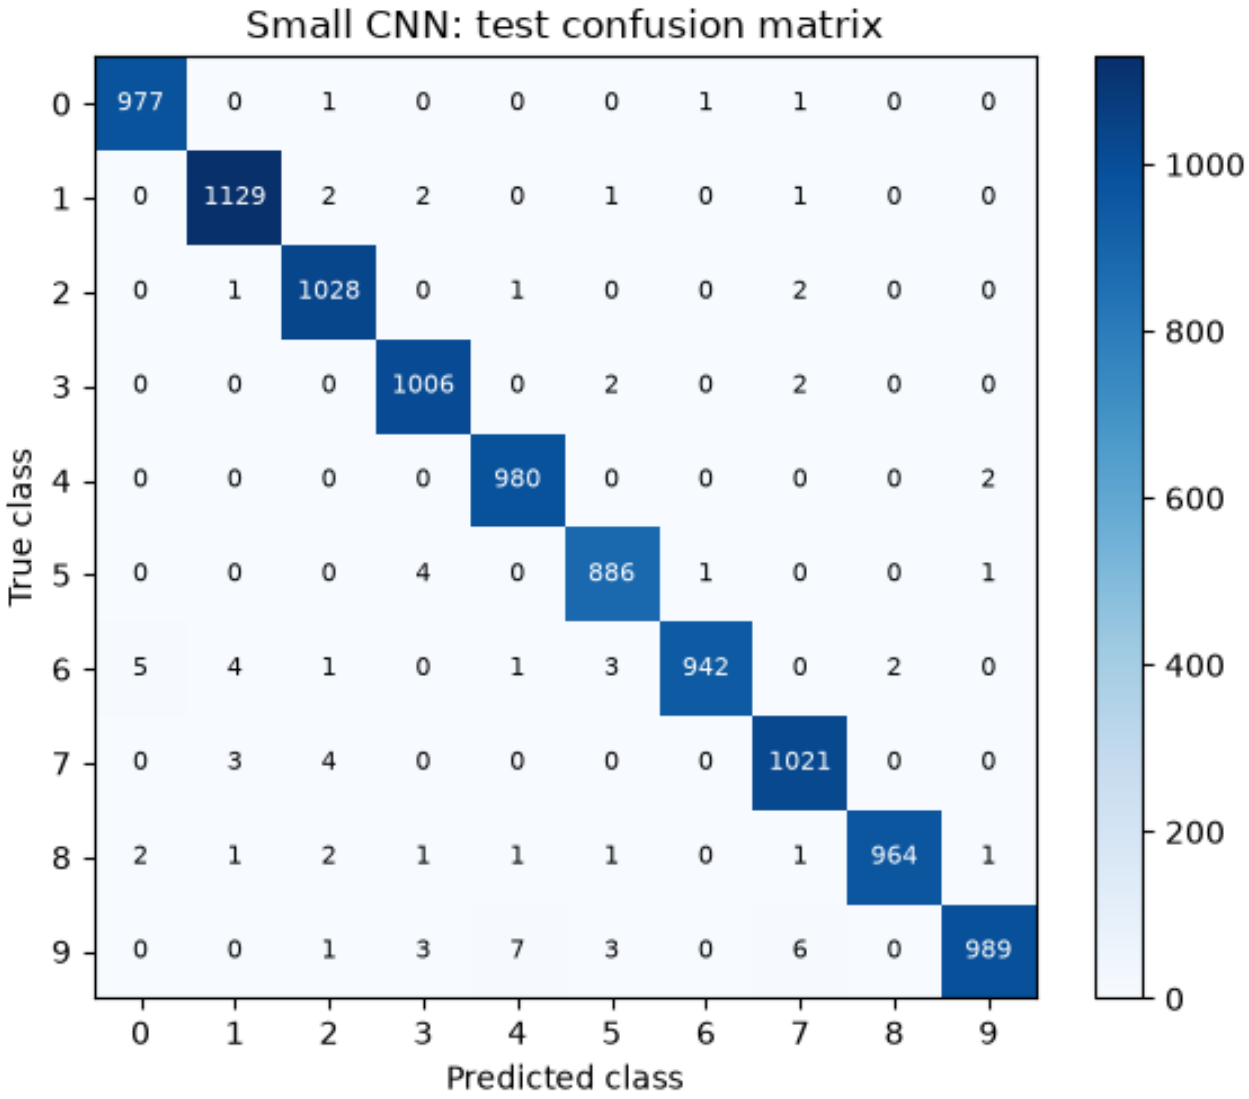





```
Small CNN
Epoch 1: train 98.33%, eval 98.48%
Epoch 2: train 98.65%, eval 98.67%
Epoch 3: train 99.22%, eval 99.08%
Epoch 4: train 99.31%, eval 99.09%
Epoch 5: train 99.20%, eval 98.92%
Epoch 6: train 99.33%, eval 99.01%
Epoch 7: train 99.52%, eval 99.21%
Epoch 8: train 99.58%, eval 99.17%
Epoch 9: train 99.56%, eval 99.15%
Epoch 10: train 99.68%, eval 99.22%
```

Epoch	Train loss	Eval loss	Train acc	Eval acc
1	0.0547	0.0514	98.33	98.48
2	0.0426	0.0416	98.65	98.67
3	0.0267	0.0288	99.22	99.08
4	0.0227	0.0282	99.31	99.09
5	0.0258	0.0343	99.20	98.92
6	0.0221	0.0320	99.33	99.01
7	0.0162	0.0238	99.52	99.21
8	0.0144	0.0248	99.58	99.17
9	0.0148	0.0263	99.56	99.15
10	0.0110	0.0254	99.68	99.22






```
print("Model       Parameters  Train acc  Test acc  Test loss  Time (s)")
for name, result in mnist_results.items():
    h = result["history"][-1]
    print(f"{name:11s} {result['parameters']:10d} {h['train_acc']:10.2f}"
          f" {h['test_acc']:9.2f} {h['test_loss']:10.4f} {result['seconds']:9.1f}")

a = mnist_results["LeNet-5"]["history"][-1]
b = mnist_results["Small CNN"]["history"][-1]

for name, result in mnist_results.items():
    matrix = np.array(result["confusion"])
    off_diagonal = matrix.copy()
    np.fill_diagonal(off_diagonal, 0)
    true, pred = np.unravel_index(off_diagonal.argmax(), off_diagonal.shape)
```

```
model = mnist_models["Small CNN"]
_, _, labels, predictions = evaluate(model, mnist_test_loader)
wrong = np.flatnonzero(labels != predictions)[:10]
fig, axes = plt.subplots(2, 5, figsize=(10, 4))
for ax in axes.flat:
    ax.axis("off")
for ax, i in zip(axes.flat, wrong):
    ax.imshow(mnist_test.data[i], cmap="gray")
    ax.set_title(f"True: {labels[i]}, pred: {predictions[i]}", fontsize=9)
fig.suptitle("Small CNN: first ten test errors")
save_figure("mnist_errors")
```

```
Model       Parameters  Train acc  Test acc  Test loss  Time (s)
LeNet-5          61706      99.72     98.95     0.0324      50.7
Small CNN        29290      99.68     99.22     0.0254     151.6
```

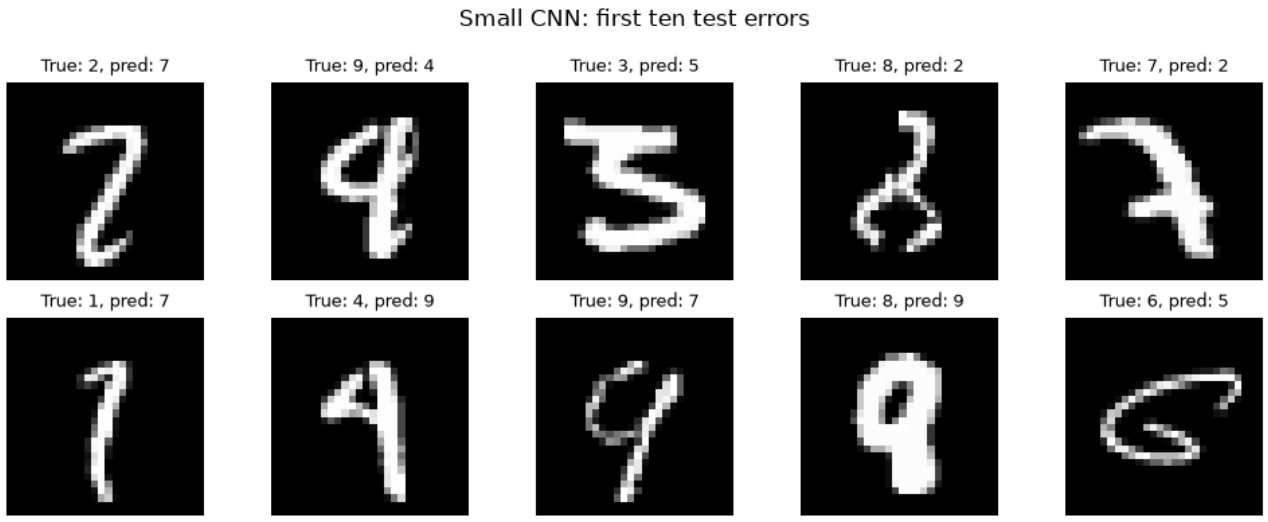
Comparison and discussion

LeNet-5 achieves **98.95%** test accuracy and the small CNN achieves **99.22%**. The CNN minus LeNet-5 difference is **0.27 percentage points**. Their training-minus-test gaps are 0.77 and 0.46 points, respectively. A smaller gap is useful, but should be considered together with the absolute test accuracy and loss.

Both models use spatial filters and shared weights, which suit handwritten digits. The modern CNN combines batch normalization, ReLU and dropout. This comparison tests the complete architectures; it cannot isolate the contribution of any one component. Cross-entropy also depends on confidence, so it can worsen even when accuracy improves if the remaining mistakes become more confident. These are single-seed results, not evidence that a small accuracy difference will hold across repeated runs.

For **LeNet-5**, the most frequent directed confusion is true digit **7** predicted as **2**, with **13** images. The off-diagonal entries count mistakes; the diagonal counts correct predictions.

For **Small CNN**, the most frequent directed confusion is true digit **9** predicted as **4**, with **7** images. The off-diagonal entries count mistakes; the diagonal counts correct predictions.



Reading the MNIST results

The small CNN has fewer parameters but takes longer in this run. Parameter count measures stored weights, while computation also depends on the number of spatial positions processed by each convolution; batch normalization adds work too. The displayed errors include 4/9 and 3/5 confusions with overlapping stroke patterns. These examples help explain some mistakes, but do not explain every test error.

3. Transfer learning with ResNet18

The ImageNet-1K pretrained ResNet18 is adapted to ants and bees using the hymenoptera dataset from the PyTorch tutorial. Both methods use the same initial classifier, 10 epochs, batch size 16, SGD with learning rate

0.001 and momentum 0.9, and a StepLR schedule (multiply the rate by 0.1 after 7 epochs).

Fine-tuning updates every layer. Feature extraction updates only the new two-class linear layer; the backbone and its batch-normalization statistics remain fixed. Training uses random resized crops to 224x224 and horizontal flips. Evaluation uses resize to 256 followed by a 224x224 center crop. Both use ImageNet normalization. The supplied `val` split is a validation set, not an independent test set. Results below use the final epoch for both methods, with no best-epoch selection.

```
from torchvision.datasets.utils import download_and_extract_archive

data_dir = Path("data/hymenoptera_data")
if not data_dir.exists():
    download_and_extract_archive(
        "https://download.pytorch.org/tutorial/hymenoptera_data.zip", "data")
mean, std = [0.485, 0.456, 0.406], [0.229, 0.224, 0.225]
train_transform = transforms.Compose([
    transforms.RandomResizedCrop(224), transforms.RandomHorizontalFlip(),
    transforms.ToTensor(), transforms.Normalize(mean, std)])
eval_transform = transforms.Compose([
    transforms.Resize(256), transforms.CenterCrop(224),
    transforms.ToTensor(), transforms.Normalize(mean, std)])
insect_train = datasets.ImageFolder(data_dir / "train", train_transform)
insect_train_eval = datasets.ImageFolder(data_dir / "train", eval_transform)
insect_val = datasets.ImageFolder(data_dir / "val", eval_transform)
insect_train_eval_loader = make_loader(insect_train_eval, 16)
insect_val_loader = make_loader(insect_val, 16)
print("Classes:", insect_train.classes)
print("Training / validation images:", len(insect_train), len(insect_val))
for split, dataset in [("train", insect_train), ("val", insect_val)]:
    print(split, dict(zip(dataset.classes, np.bincount(dataset.targets).tolist())))

transfer_results = {}
transfer_models = {}
for name, frozen in [("Fine-tuning", False), ("Feature extraction", True)]:
    set_seed()
    model = models.resnet18(weights=models.ResNet18_Weights.IMAGENET1K_V1)
    if frozen:
        for parameter in model.parameters():
            parameter.requires_grad = False
    model.fc = nn.Linear(model.fc.in_features, 2)
    model = model.to(device)
    optimizer = torch.optim.SGD(
        [p for p in model.parameters() if p.requires_grad], lr=0.001, momentum=0.9)
    scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=7, gamma=0.1)
    train_loader = make_loader(insect_train, 16, True)
    print("\n" + name, flush=True)

    history, seconds = train_model(model, train_loader, insect_train_eval_loader,
                                   insect_val_loader, optimizer, scheduler=scheduler,
                                   frozen_backbone=frozen)
    _, _, labels, predictions = evaluate(model, insect_val_loader)
    transfer_results[name] = dict(history=history, seconds=seconds,
        parameters=sum(p.numel() for p in model.parameters() if p.requires_grad))
    transfer_models[name] = model
    print_history(history)
    plot_history(history, name + ": ants and bees", name.replace(" ", "_") + "_curves",
"Validation")
    matrix = confusion_plot(labels, predictions, insect_train.classes,
                            name + ": validation confusion", name.replace(" ", "_") + "_confusion")
    transfer_results[name]["confusion"] = matrix.tolist()
```

```
Classes: ['ants', 'bees']
Training / validation images: 244 153
train {'ants': 123, 'bees': 121}
val {'ants': 70, 'bees': 83}

Fine-tuning

Epoch 1: train 90.57%, eval 90.20%

Epoch 2: train 95.49%, eval 94.12%

Epoch 3: train 97.95%, eval 94.77%
```

```
Epoch 4: train 98.36%, eval 94.12%

Epoch 5: train 98.77%, eval 94.12%

Epoch 6: train 99.18%, eval 93.46%

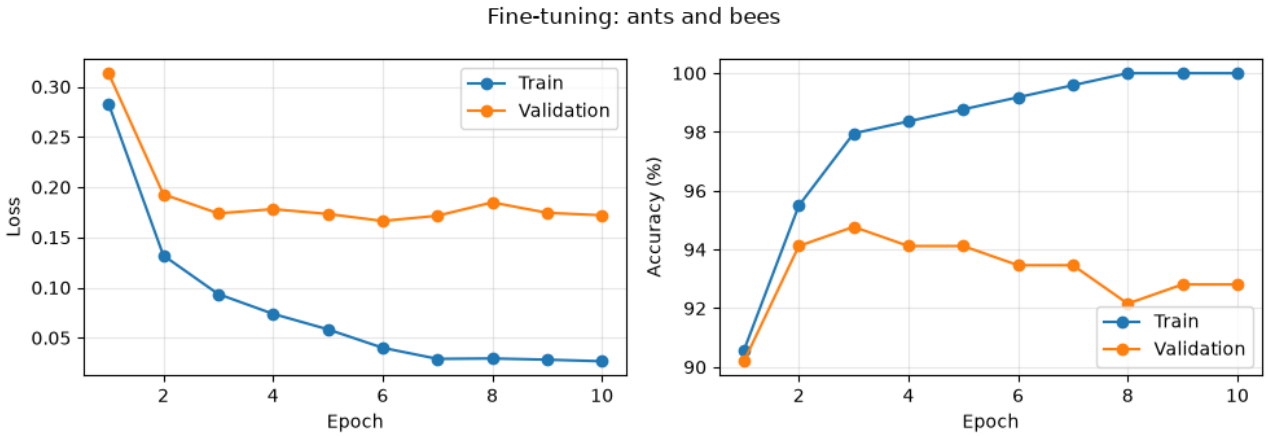
Epoch 7: train 99.59%, eval 93.46%

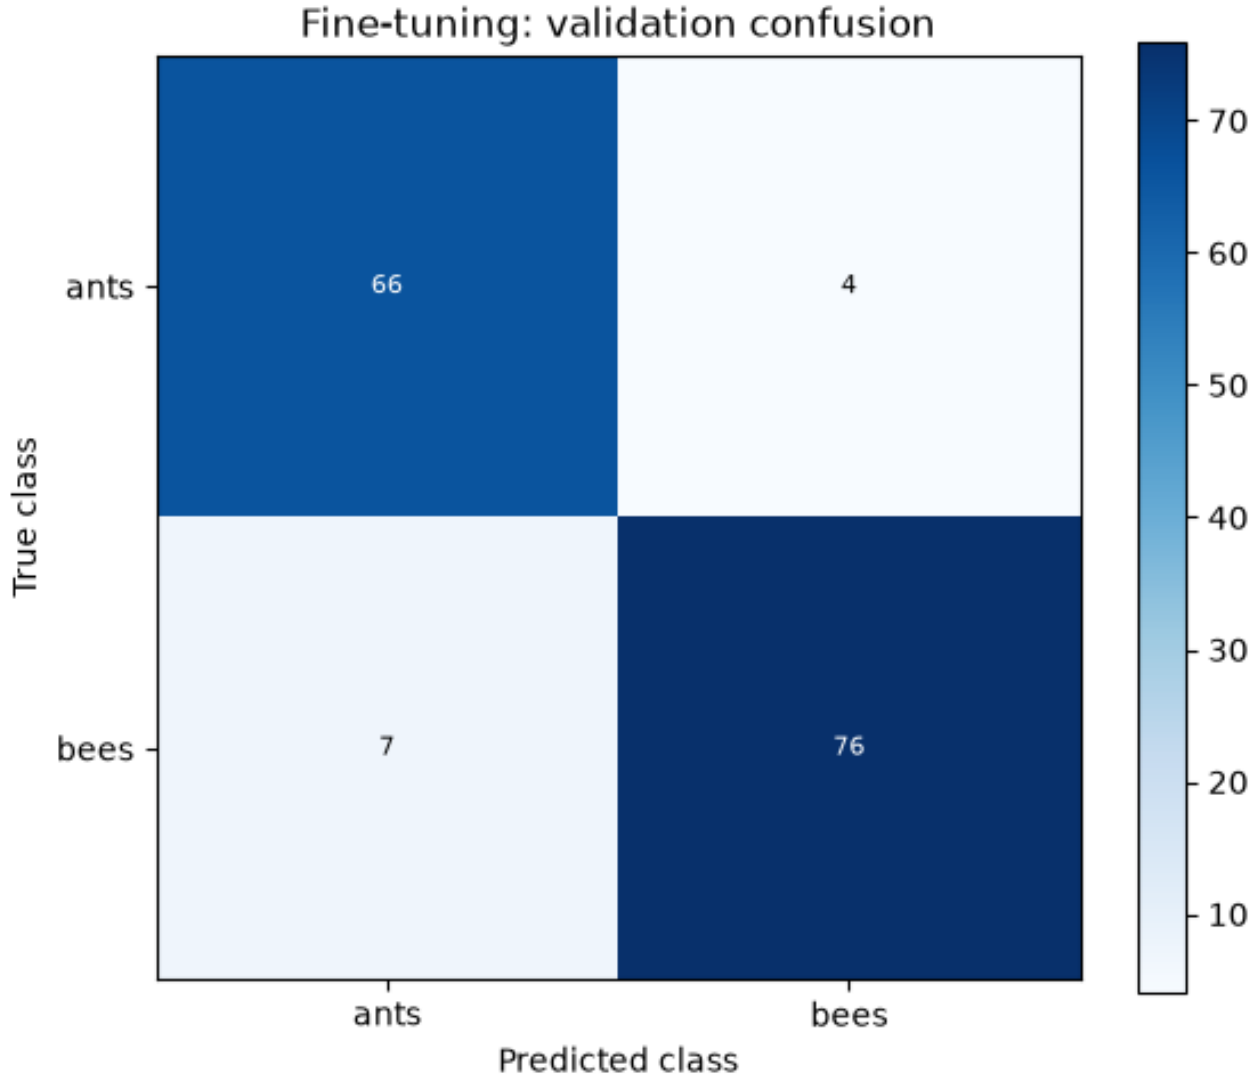
Epoch 8: train 100.00%, eval 92.16%

Epoch 9: train 100.00%, eval 92.81%

Epoch 10: train 100.00%, eval 92.81%
```

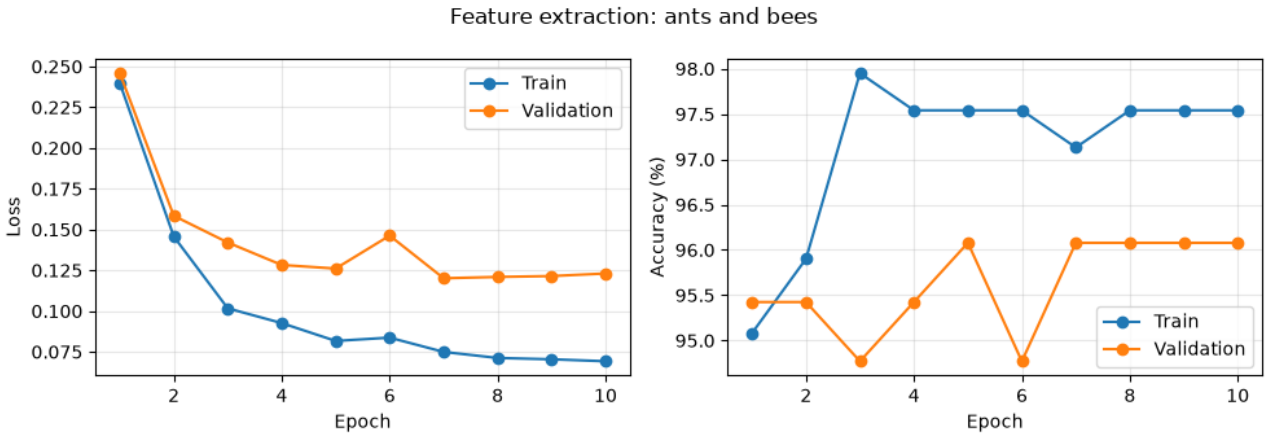
Epoch	Train loss	Eval loss	Train acc	Eval acc
1	0.2832	0.3134	90.57	90.20
2	0.1320	0.1929	95.49	94.12
3	0.0936	0.1739	97.95	94.77
4	0.0739	0.1781	98.36	94.12
5	0.0584	0.1733	98.77	94.12
6	0.0401	0.1663	99.18	93.46
7	0.0290	0.1715	99.59	93.46
8	0.0294	0.1848	100.00	92.16
9	0.0282	0.1745	100.00	92.81
10	0.0266	0.1720	100.00	92.81

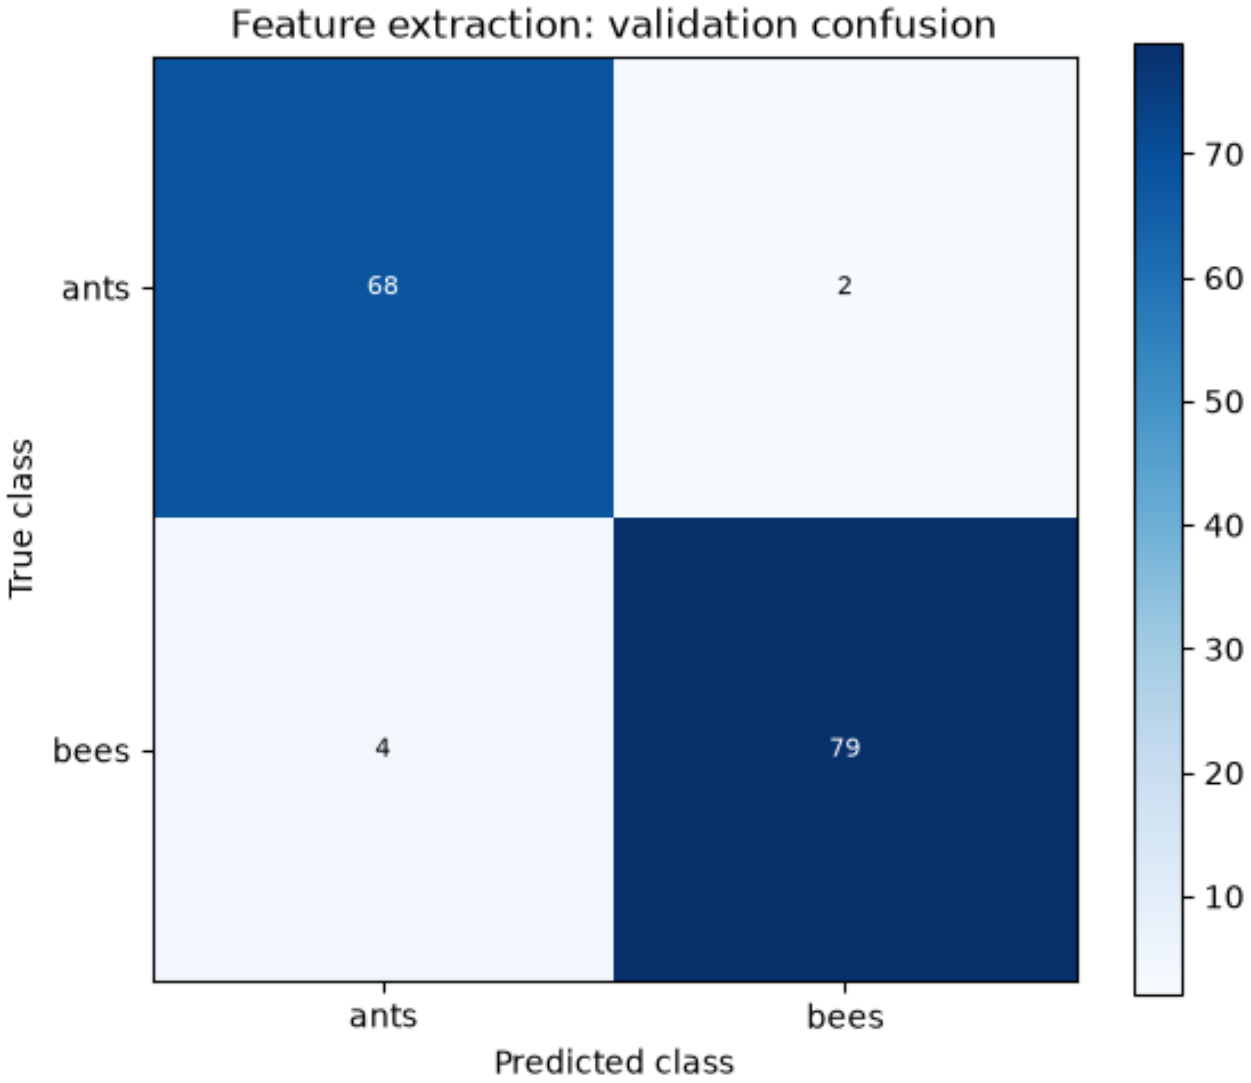





```
Feature extraction
Epoch 1: train 95.08%, eval 95.42%
Epoch 2: train 95.90%, eval 95.42%
Epoch 3: train 97.95%, eval 94.77%
Epoch 4: train 97.54%, eval 95.42%
Epoch 5: train 97.54%, eval 96.08%
Epoch 6: train 97.54%, eval 94.77%
Epoch 7: train 97.13%, eval 96.08%
Epoch 8: train 97.54%, eval 96.08%
Epoch 9: train 97.54%, eval 96.08%
Epoch 10: train 97.54%, eval 96.08%
```

Epoch	Train loss	Eval loss	Train acc	Eval acc
1	0.2399	0.2459	95.08	95.42
2	0.1457	0.1584	95.90	95.42
3	0.1017	0.1421	97.95	94.77
4	0.0926	0.1283	97.54	95.42
5	0.0817	0.1261	97.54	96.08
6	0.0837	0.1463	97.54	94.77
7	0.0749	0.1201	97.13	96.08
8	0.0712	0.1209	97.54	96.08
9	0.0704	0.1215	97.54	96.08
10	0.0692	0.1230	97.54	96.08






```
print("Method              Trainable params  Train acc  Val acc  Val loss  Time (s)")
for name, result in transfer_results.items():
    h = result["history"][-1]
    print(f"{name:19s} {result['parameters']:16d} {h['train_acc']:10.2f}"
          f" {h['test_acc']:8.2f} {h['test_loss']:9.4f} {result['seconds']:9.1f}")

fig, axes = plt.subplots(2, 4, figsize=(10, 5))
indices = [insect_val.targets.index(0), insect_val.targets.index(1)]
indices += [i for i in range(len(insect_val)) if i not in indices][:2]
for row, (name, model) in enumerate(transfer_models.items()):
    model.eval()
    for col, i in enumerate(indices):
        x, y = insect_val[i]
        with torch.no_grad():
            pred = model(x.unsqueeze(0).to(device)).argmax(1).item()
```

```
        image = (x * torch.tensor(std)[:, None, None] + torch.tensor(mean)[:, None, None])
        axes[row, col].imshow(image.permute(1, 2, 0).clamp(0, 1))
        axes[row, col].set_title(name + "\nTrue: " + insect_val.classes[y] +
                                 ", pred: " + insect_val.classes[pred], fontsize=8)
        axes[row, col].axis("off")
save_figure("transfer_predictions")

fine = transfer_results["Fine-tuning"]["history"][-1]
fixed = transfer_results["Feature extraction"]["history"][-1]

results = dict(seed=SEED, device=str(device), torch_version=str(torch.__version__),
               dense=dict(history=dense_history, seconds=dense_seconds),
               mnist=mnist_results, transfer=transfer_results)
Path("results/metrics.json").write_text(json.dumps(results, indent=2))
print("Saved measured results to results/metrics.json")
```

```
Method              Trainable params  Train acc  Val acc  Val loss  Time (s)
Fine-tuning                 11177538     100.00    92.81    0.1720     190.9
Feature extraction              1026      97.54    96.08    0.1230     117.8
```



Comparison and discussion

Fine-tuning reaches **92.81%** validation accuracy with loss **0.1720**. Feature extraction reaches **96.08%** with loss **0.1230**. The accuracy difference (fine-tuning minus feature extraction) is **-3.27 percentage points**. The corresponding training accuracies are 100.00% and 97.54%. Training images are evaluated with center crops for these figures.

Fine-tuning can adjust the pretrained visual features to this dataset, but fits many more parameters from only 244 training images. Feature extraction relies on the features already learned from ImageNet and trains only the classifier. The timing table shows the computational cost of each approach on this machine; it includes training and per-epoch evaluation, but excludes downloading weights. The fixed backbone still needs a forward pass for every image.

There are only 153 validation images, so one additional correct prediction changes accuracy by 0.65 percentage points. Small differences should therefore be interpreted cautiously. One seed and one split do not establish a general ranking, and no separate held-out test set is available here. Freezing the backbone's batch-normalization statistics makes this a strict feature extractor; the tutorial's training-mode implementation can update those statistics.

```
Saved measured results to results/metrics.json
```

Interpreting the transfer-learning curves

Fine-tuning reaches 100.00% training accuracy, but its final validation accuracy is 92.81%, below its highest observed value of 94.77%. This pattern is consistent with overfitting: a better fit to the training images does not guarantee better validation accuracy. Feature extraction ends at 96.08% validation accuracy while updating far fewer parameters. For this run, keeping the pretrained features fixed works better at the final epoch. This does not mean fine-tuning is generally worse; its result depends on the learning rate, training duration and amount of data. The reported comparison uses epoch 10 for both methods.